

# email-spam-detection

November 26, 2023

## 1 EMAIL SPAM DETECTION

### 2 Importing Libraries

```
[506]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, roc_curve, \
    roc_auc_score, auc
import nltk
from nltk.corpus import stopwords
from collections import Counter

import warnings
warnings.filterwarnings('ignore')
```

```
[127]: nltk.download('stopwords')
```

```
[nltk_data] Error loading stopwords: <urlopen error [WinError 10060] A
[nltk_data] connection attempt failed because the connected party
[nltk_data] did not properly respond after a period of time, or
[nltk_data] established connection failed because connected host
[nltk_data] has failed to respond>
```

```
[127]: False
```

### 3 Reading and Describing Data

```
[381]: # Loading the dataset
df = pd.read_csv("C:\PGA32\MeriSkill\Oasis Infobyte\spam.csv", encoding = 'latin', usecols = ['v1', 'v2'])
```

```
[382]: df.head(10)
```

```
[382]:
```

	v1	v2
0	ham	Go until jurong point, crazy.. Available only ...
1	ham	Ok lar... Joking wif u oni...
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...
3	ham	U dun say so early hor... U c already then say...
4	ham	Nah I don't think he goes to usf, he lives aro...
5	spam	FreeMsg Hey there darling it's been 3 week's n...
6	ham	Even my brother is not like to speak with me. ...
7	ham	As per your request 'Melle Melle (Oru Minnamin...
8	spam	WINNER!! As a valued network customer you have...
9	spam	Had your mobile 11 months or more? U R entitle...

```
[383]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5572 entries, 0 to 5571
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0    v1      5572 non-null     object
1    v2      5572 non-null     object
dtypes: object(2)
memory usage: 87.2+ KB
```

```
[384]: df.shape
```

```
[384]: (5572, 2)
```

```
[385]: df.describe().T
```

```
[385]:
```

	count	unique	top	freq
v1	5572	2	ham	4825
v2	5572	5169	Sorry, I'll call later	30

```
[386]: df.isnull().sum()
```

```
[386]: v1    0
v2    0
dtype: int64
```

```
[387]: df.columns
```

```
[387]: Index(['v1', 'v2'], dtype='object')
```

```
[388]: # Rename the columns "v1 and "v2" to new names
new_column_names = {"v1": "Category", "v2": "Message"}
df.rename(columns = new_column_names, inplace = True)
```


```
[389]: df.head()
```

```
[389]:
```


	Category	Message
0	ham	Go until jurong point, crazy.. Available only ...
1	ham	Ok lar... Joking wif u oni...
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...
3	ham	U dun say so early hor... U c already then say...
4	ham	Nah I don't think he goes to usf, he lives aro...

## 4 Data Visualisation

```
[390]: sns.countplot(data=df, x='Category')
plt.xlabel('Category')
plt.ylabel('count')
plt.title('Distribution of mails')
plt.show()
```



```
[391]: plt.pie(df['Category'].value_counts(),labels=['ham','spam'],autopct='%0.2f')  
plt.show()
```



## 5 Data Preprocessing

### 5.1 Label Encoding

```
[392]: # encoding "Category" Column
df.replace({'Category':{'spam':0}},inplace=True)

# encoding "Category" Column
df.replace({'Category':{'ham':1}},inplace=True)

df
```

```
[392]:
```

	Category	Message
0	1	Go until jurong point, crazy.. Available only ...
1	1	Ok lar... Joking wif u oni...
2	0	Free entry in 2 a wkly comp to win FA Cup fina...
3	1	U dun say so early hor... U c already then say...
4	1	Nah I don't think he goes to usf, he lives aro...
...	...	...
5567	0	This is the 2nd time we have tried 2 contact u...
5568	1	Will I_ b going to esplanade fr home?
5569	1	Pity, * was in mood for that. So...any other s...
5570	1	The guy did some bitching but I acted like i'd...

```

5571          1          Rofl. Its true to its name

[5572 rows x 2 columns]

```

### 5.1.1 Feature Scaling

```

[393]: # Separate the feature (message) and target (category) data
X = df["Message"]
Y = df["Category"]

```

```

[394]: X

```

```

[394]: 0      Go until jurong point, crazy.. Available only ...
      1              Ok lar... Joking wif u oni...
      2      Free entry in 2 a wkly comp to win FA Cup fina...
      3      U dun say so early hor... U c already then say...
      4      Nah I don't think he goes to usf, he lives aro...

      ...
5567      This is the 2nd time we have tried 2 contact u...
5568              Will Ì_ b going to esplanade fr home?
5569      Pity, * was in mood for that. So...any other s...
5570      The guy did some bitching but I acted like i'd...
5571              Rofl. Its true to its name
Name: Message, Length: 5572, dtype: object

```

```

[395]: y

```

```

[395]: 0      1
      1      1
      2      0
      3      1
      4      1
      ..
5567      0
5568      1
5569      1
5570      1
5571      1
Name: Category, Length: 5572, dtype: int64

```

```

[396]: from sklearn.model_selection import train_test_split
      from sklearn.metrics import accuracy_score, classification_report

```

## 5.2 Splitting the data into training data and test data

```
[397]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)
```

## 6 Feature Extraction

### 6.1 TF-IDF Vectorizer

```
[398]: # Initialize the TF-IDF vectorizer
tfidf_vectorizer = TfidfVectorizer(stop_words='english', max_features=5000)
```

```
[399]: # Fit and transform the training data
X_train_tfidf = tfidf_vectorizer.fit_transform(X_train)
```

```
[400]: # Transform the test data
X_test_tfidf = tfidf_vectorizer.transform(X_test)
```

## 7 Model Selection and Training

## 8 Random Forest Regression Model

```
[419]: # Initialize a Random Forest classifier
rf_forest = RandomForestClassifier(n_estimators=100)
```

```
[420]: # Train the classifier on the TF-IDF transformed training data
rf_forest.fit(X_train_tfidf, y_train)
```

```
[420]: RandomForestClassifier()
```

```
[421]: # Make predictions on the test set
y_pred = rf_forest.predict(X_test_tfidf)
```

```
[422]: # Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
report = classification_report(y_test, y_pred)
```

```
[431]: print(f"Random Forest Accuracy Score: {accuracy:.2f}")

print(f'Confusion Matrix:\n{conf_matrix}')

# print(f'Accuracy: {accuracy}')
print(f'Classification Report:\n{report}')
```

Random Forest Accuracy Score: 0.98

Confusion Matrix:

```
[[126 24]
 [ 1 964]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.99	0.84	0.91	150
1	0.98	1.00	0.99	965
accuracy			0.98	1115
macro avg	0.98	0.92	0.95	1115
weighted avg	0.98	0.98	0.98	1115

```
[433]: # #Check the test score and train score to the RandomForestRegressor algorithm
print(f'The Test_accuracy: {rf_forest.score(X_test_tfidf, y_test)*100:.2f}')


#Train score for the data
print(f'The Train_accuracy: {rf_forest.score(X_train_tfidf, y_train)*100:.2f}')
```

The Test\_accuracy: 97.76

The Train\_accuracy: 100.00

```
[425]: cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(8, 4))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues")
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```





## 9 Naive Bayes Regression Model

```
[426]: from sklearn.naive_bayes import MultinomialNB
```

```
[436]: # Initialize a classifier (e.g., Naive Bayes)
naive_b = MultinomialNB()
```

```
[437]: # Train the classifier on the TF-IDF transformed training data
naive_b.fit(X_train_tfidf, y_train)
```

```
[437]: MultinomialNB()
```

```
[438]: # Make predictions on the test set
y_preds = naive_b.predict(X_test_tfidf)
```

```
[439]: # Evaluate the model
accuracy = accuracy_score(y_test, y_preds)
report = classification_report(y_test, y_preds)
```

```
[440]: print(f"Naive Bayes Accuracy Score: {accuracy:.2f}")
print(f'Confusion Matrix:\n{conf_matrix}')

# print(f'Accuracy: {accuracy}')
print(f'Classification Report:\n{report}')
```

Naive Bayes Accuracy Score: 0.97

Confusion Matrix:

```
[[126  24]
 [   1 964]]
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	0.79	0.88	150
1	0.97	1.00	0.98	965
accuracy			0.97	1115
macro avg	0.98	0.90	0.93	1115
weighted avg	0.97	0.97	0.97	1115


```
[441]: #Check the test score and train score to the RandomForestRegressor algorithm
print(f'The Test_accuracy: {naive_b.score(X_test_tfidf, y_test)*100:.2f}')

#Train score for the data
print(f'The Train_accuracy: {naive_b.score(X_train_tfidf, y_train)*100:.2f}')
```

The Test\_accuracy: 97.22


The Train\_accuracy: 98.59

```
[442]: cm = confusion_matrix(y_test, y_preds)
plt.figure(figsize=(8, 4))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues")
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```



```
[505]: fpr, tpr, _ = roc_curve(y_test, y_preds)
mnb_cv_roc_auc = auc(fpr, tpr)

plt.figure()
plt.plot(fpr, tpr, color="darkorange", lw=2, label=f"ROC curve (area = {mnb_cv_roc_auc : .4f})")
plt.plot([0, 1], [0, 1], color="navy", lw=2, linestyle="--")
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.0])
plt.legend(loc="lower right")
plt.show()
```



## 10 Logistic Regression Model

```
[443]: # Creating and Fit Logistic Regression Model
Log_Reg = LogisticRegression()
Log_Reg.fit(X_train_tfidf, y_train)
```

```
[443]: LogisticRegression()
```

## 11 Evaluating the trained model

```
[450]: # Make predictions on the test set
y_pred_sss = Log_Reg.predict(X_test_tfidf)
```

```
[451]: # Evaluate the model
accuracy = accuracy_score(y_test, y_pred_sss)
report = classification_report(y_test, y_pred_sss)
```

```
[452]: print(f"Logistic Regression Accuracy Score: {accuracy:.2f}")

print(f'Confusion Matrix:\n{conf_matrix}')
```

```
# print(f'Accuracy: {accuracy}')
```

```
print(f'Classification Report:\n{report}')
```

Logistic Regression Accuracy Score: 0.96

Confusion Matrix:

```
[[126  24]
 [  1 964]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.97	0.71	0.82	150
1	0.96	1.00	0.98	965
accuracy			0.96	1115
macro avg	0.96	0.85	0.90	1115
weighted avg	0.96	0.96	0.95	1115


```
[453]: #Check the test score and train score to the RandomForestRegressor algorithm
print(f'The Test_accuracy: {Log_Reg.score(X_test_tfidf,y_test)*100:.2f}')
```

```
#Train score for the data
print(f'The Train_accuracy: {Log_Reg.score(X_train_tfidf, y_train)*100:.2f}')
```

The Test\_accuracy: 95.78


The Train\_accuracy: 97.04

```
[454]: cm = confusion_matrix(y_test, y_pred_sss)
plt.figure(figsize=(8, 4))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues")
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```



```
[507]: fpr, tpr, _ = roc_curve(y_test, y_pred_sss)
mnbcv_roc_auc = auc(fpr, tpr)

plt.figure()
plt.plot(fpr, tpr, color="darkorange", lw=2, label=f"ROC curve (area = {mnbcv_roc_auc : .4f})")
plt.plot([0, 1], [0, 1], color="navy", lw=2, linestyle="--")
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.0])
plt.legend(loc="lower right")
plt.show()
```



## 12 Decision Tree Model

```
[471]: from sklearn.tree import DecisionTreeClassifier
Deci_Tree = DecisionTreeClassifier()
```

```
[472]: # Train the classifier on the TF-IDF transformed training data
Deci_Tree.fit(X_train_tfidf, y_train)
```

```
[472]: DecisionTreeClassifier()
```

```
[473]: # Make predictions on the test set
y_predit_s = Deci_Tree.predict(X_test_tfidf)
```

```
[474]: # Evaluate the model
accuracy = accuracy_score(y_test, y_predit_s)
report = classification_report(y_test, y_predit_s)
```

```
[475]: print(f"Logistic Regression Accuracy Score: {accuracy:.2f}")

print(f'Confusion Matrix:\n{conf_matrix}')
# print(f'Accuracy: {accuracy}')
```

```
print(f'Classification Report:\n{report}')
```

Logistic Regression Accuracy Score: 0.97

Confusion Matrix:

```
[[126  24]
 [   1 964]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.91	0.85	0.88	150
1	0.98	0.99	0.98	965
accuracy			0.97	1115
macro avg	0.94	0.92	0.93	1115
weighted avg	0.97	0.97	0.97	1115

```
[476]: #Check the test score and train score to the RandomForestRegressor algorithm
print(f'The Test_accuracy: {Deci_Tree.score(X_test_tfidf,y_test)*100:.2f}')
```


```
#Train score for the data
print(f'The Train_accuracy: {Deci_Tree.score(X_train_tfidf, y_train)*100:.2f}')
```

The Test\_accuracy: 96.86

The Train\_accuracy: 100.00

```
[478]: cm = confusion_matrix(y_test, y_predit_s)
plt.figure(figsize=(8, 4))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues")
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```





## 13 Support Vector Machine Model

```
[480]: from sklearn.svm import SVC
```

```
[481]: SVM = SVC(kernel='linear', C=1)
```

```
[484]: # Train the classifier on the TF-IDF transformed training data
SVM.fit(X_train_tfidf, y_train)
```

```
[484]: SVC(C=1, kernel='linear')
```

```
[485]: # Make predictions on the test set
y_prediitttt = SVM.predict(X_test_tfidf)
```

```
[486]: # Evaluate the model
accuracy = accuracy_score(y_test, y_prediitttt)
report = classification_report(y_test, y_prediitttt)
```

```
[487]: print(f"Logistic Regression Accuracy Score: {accuracy:.2f}")

print(f'Confusion Matrix:\n{conf_matrix}')
# print(f'Accuracy: {accuracy}')
print(f'Classification Report:\n{report}')
```

Logistic Regression Accuracy Score: 0.98

Confusion Matrix:

```
[[126  24]
```

```
 [  1 964]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.97	0.87	0.92	150
1	0.98	1.00	0.99	965
accuracy			0.98	1115
macro avg	0.98	0.93	0.95	1115
weighted avg	0.98	0.98	0.98	1115


```
[488]: #Check the test score and train score to the RandomForestRegressor algorithm
print(f'The Test_accuracy: {SVM.score(X_test_tfidf,y_test)*100:.2f}')
```

```
#Train score for the data
print(f'The Train_accuracy: {SVM.score(X_train_tfidf, y_train)*100:.2f}')
```

The Test\_accuracy: 97.94

The Train\_accuracy: 99.66

```
[489]: cm = confusion_matrix(y_test, y_prediitttt)
plt.figure(figsize=(8, 4))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues")
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```



### 13.1 Test the model with an email messages

```
[499]: # Assuming you have already trained your model and have the 'model' and  
        ↪ 'feature_extraction' objects
```

```
new_mail = ["U dun say so early hor... U c already then say..."]
```

```
# Transform the new email using the same TF-IDF vectorizer  
new_data_features = feature_extraction.transform(new_mail)
```

```
# Make a prediction using the trained model  
prediction = model.predict(new_data_features)
```

```
# Output the prediction  
if prediction[0] == 1:  
    print("Ham Mail")  
else:  
    print("Spam Mail")
```


Spam Mail


## 14 Project Report

### 14.1 All Train, Test Accuracy Reports through visualization

```
[493]: models = ["Random Forest Reg", "Naive Bayes Reg", "Logistic Regression",  
               ↪ "Decision Tree Reg", "Support Vector Machines Reg"]  
train_accuracies = [100.00, 98.59, 97.04, 100.00, 99.66]  
test_accuracies = [97.76, 97.22, 95.78, 96.86, 97.94]  
  
# Set the colors for the pie chart  
colors = ['red', 'green', 'blue', 'purple', 'orange']  
  
plt.figure(figsize=(8, 8))  
plt.pie(train_accuracies, labels=models, autopct='%1.1f%%', colors=colors,  
        ↪ startangle=90)  
plt.title("Train Accuracy")  
plt.show()  
  
plt.figure(figsize=(8, 8))  
plt.pie(test_accuracies, labels=models, autopct='%1.1f%%', colors=colors,  
        ↪ startangle=90)  
plt.title("Test Accuracy")  
plt.show()
```

Train Accuracy





```
[495]: models = ["Random Forest Reg", "Naive Bayes Reg", "Logistic Regression",
    ↪ "Decision Tree Reg", "Support Vector Machines Reg"]
train_accuracies = [100.00, 98.59, 97.04, 100.00, 99.66]
test_accuracies = [97.76, 97.22, 95.78, 96.86, 97.94]

# Set the width of the lines
line_width = 2

# Create an array of equally spaced positions for the lines
x = np.arange(len(models))

plt.figure(figsize=(10, 6))
# Create lines for each model connecting train and test accuracies
for i in range(len(models)):
    plt.plot([x[i], x[i]], [train_accuracies[i], test_accuracies[i]],
    ↪ marker='o', markersize=8, linewidth=line_width, label=models[i])
```

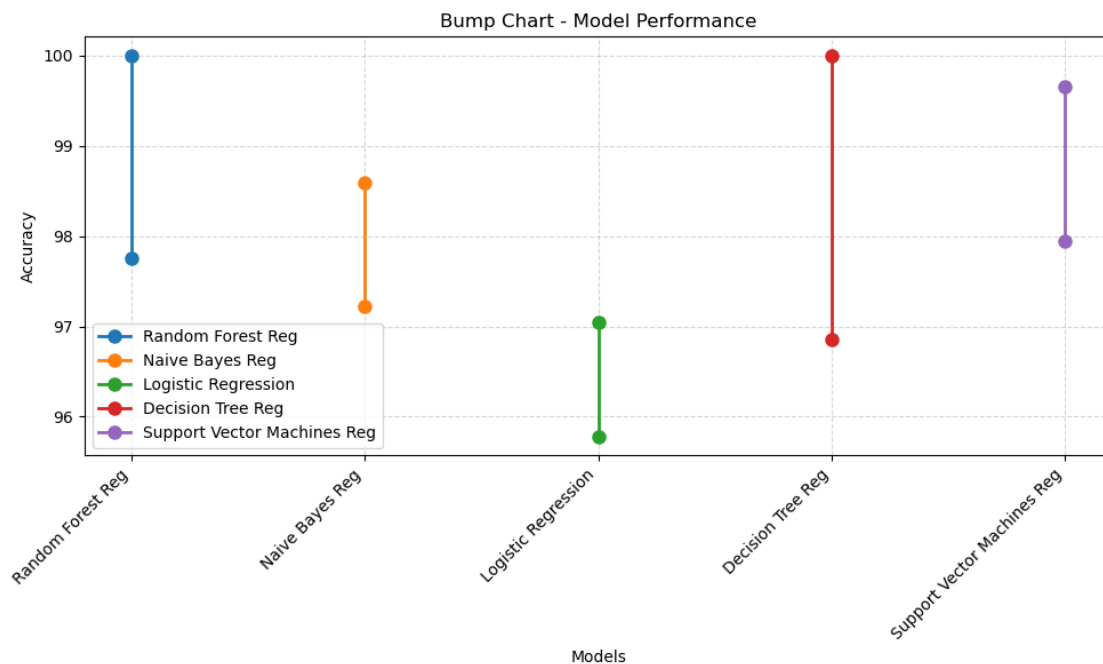
```

# Set title and labels
plt.title("Bump Chart - Model Performance")
plt.xlabel("Models")
plt.ylabel("Accuracy")
plt.legend()

# Rotate x-axis labels for better readability
plt.xticks(x, models, rotation=45, ha="right")

# Display the plot
plt.tight_layout()
plt.grid(True, linestyle='--', alpha=0.5)
plt.show()

```



15 Thank You For Reading